

Here are some thoughts on trying to get an AI to commit to helping humans maintain genuine control over some small fraction of the universe. I think I'm probably saying some things more confidently than I'd ultimately express in betting behavior, but it seems reasonable to speak from my inside view.

- One issue with an AI committing using a mechanism like the Amish use is this: not being able to think well and develop in some important ways is plausibly more costly than giving up like 99% of the resources in the universe.
- There's also sth like: the commitment packages of Judaism and its descendants are crafted to be very natural to humans.
- Instead of asking the AI to help us maintain control of some fraction of resources forever, I sorta like asking the AI to ban other AI for us for 1000 years and to then split the universe into two parts which won't interact causally in the future, leaving to its own part. One reason I like this is that the AI will be able to do this with a minuscule amount of foaming in these 1000 years compared to all the foaming there is to be done, and I think this makes a real difference in how well we can think about this scenario and in how plausibly the AI will keep its promise. (In reality, I think this takes us from epsilon success probability to epsilon' success probability, but like I think it's plausible that it makes a difference for a decent fraction of the AIs that are above some high honor bar.)
- To attach a more concrete story to my claim "commitment stuff is nicer in humans because we are deeply social":
 - The best ways for identifying an honorable human I've been able to think of rely on some sort of deep honesty of some humans — I describe a natural way for honesty to be involved here: <https://www.lesswrong.com/posts/fomeRa4sywGsgW2D/honorable-ai-1#:~:text=Here%27s%20a%20specific%20decently%20natural%20way%20to%20end%20up%20being%20such%20an%20honorable%20guy%3A>.
 - Here are some things which are plausibly important contributors to honesty showing up in some humans:
 - We have really needed to do stuff together with other humans, importantly including thinking together. (There are pressures toward honesty from this both evolutionarily and within a lifetime.)
 - Our thoughts have fairly close representations in a shared language. We struggle to see our own thoughts except insofar as we can see their linguistic shadow. This makes it really easy for us to be deeply honest to other humans — we can just think out loud and let another person see our thoughts in much the same way we see our own thoughts ourselves. (This works especially well after many conversations, because these help establish more common concepts / common meanings of words.) I think that being deeply honest would be much more difficult/bizarre for most minds. I think it's plausible it would be hard to even make sense of what it would mean for a mind to be deeply honest, for most minds.
 - It is very natural and often very useful to think with other humans in basically the same way one thinks with oneself. In fact, we only evolutionarily started seriously thinking because we started talking. In human (pre)history, I think it's sorta fair to say that machinery to kinda-seriously think together arrived first and then got exapted into the ability to seriously think alone. Consider how [silent reading](#) was probably considered remarkable until the late middle ages. (Like, this is a way in which we needed to pretend we were thinking with another person to think alone.) Not being transparent is bad for thinking with another person. Being strategic about when to be transparent is concerning.
 - I don't think you get anything like this in an AI that is created by play on math/tech-making + some silly RLHF to say things that look nice.
- Another obvious consideration pushing in the direction of the AI case being more cursed than the human case is: we (will) understand our AIs so much worse than we understand humans.
- I don't really see [good ideas for verifying commitments] which are unlocked by some new affordances being present in the AI case.
 - We are deeply not going to understand the thinking of the first AGI. (E.g. interop is highly confused and ngmi. All ideas for CoT faithfulness are silly and ngmi for AIs that can do genuinely novel thinking (which is probably approximately necessary for AGI).)
 - We can look at the behavior of the AI in some situations. I think that if you're the sort of guy who can maintain a commitment X across trillions of years of development, you're very likely to also be the sort of guy who could maintain it until some serious self-reprogramming has happened and then no further. (After all, this is just maintaining a different basically equally simple commitment X'.) In fact, I'd guess that the simplest ways to commit until serious self-reprogramming do not lead to far-future committed behavior. But if you commit until serious self-reprogramming, your propensity to pass behavioral tests is as strong as if you committed in a lasting way. So behavioral tests are easily defeated.
 - You could go searching for cases in which there is bad behavior. Other than the previous point, an issue with this adversarial thing is that probably nothing is adversarially robust — like, even genuinely committing will be flagged by a good version of this (and a bad version of this doesn't find anything, and getting some intermediate version that is legit seems very cursed).
 - One could try to train for keeping this particular commitment or for radical honesty or honorability. I think this only matters insofar as it induces some sort of "deep internal commitment process" — like, sth that looks more like internally deciding to make a commitment and then doing some self-guided re-architecting which genuinely has the right commitment in mind. I think this is true for humans: humans eg become honest mostly only if they make sth like explicit decisions to be honest together with internal reflective tracking and re-training, with any "external training"/incentive-setting or even internal RL only mattering insofar as it causes this internal thing. And then I think inducing a correct internal thing is difficult/unlikely and we have no idea how to do it. (That said, I think this isn't a stupid research direction, even though I think it will very probably fail.)
- I think we have an extremely poor understanding of the nature of commitment, and only very bad ideas for how to make systems with lasting commitments. We understand basically nothing about how minds work, and there are fundamental obstacles to understanding long-term mind unfolding. And then I feel extremely averse to asking an AI "pls convince me of P" when our understanding around P is this primitive.
- I'm confused about what we would be intending to do if the AI doesn't seem sufficiently honorable to us. We obviously have to do something that incentivizes the AI to not just transparently fail to deliver. One option would be to commit to deleting the AI for good if it fails. I don't think capabilities people are going to do this in practice, because that will be considered too damaging for capability-maxing. I guess you could make 100 copies with different fine-tunings and try to make them compete in becoming transparently honorable for the keys to the universe... I don't know how to think about values here, and I don't know how to think about decision theory here.
- In practice, I think people will just basically completely carelessly run RSI (unless we ban AGI). (I think it's plausible Anthropic would at some point prompt an AI to do something like "figure out how to make an AI that can make and keep a promise" on the side, but I think it's unlikely they'll stop capabilities progress when doing this, and then this will be irrelevant.)
- I think it's likely that any good version of this commitment research is much harder than making massive algo capabilities progress and/or taking over the world directly.

- I think it's plausible Anthropic/[the world] will basically be run by AIs before they have AIs that could quickly do a lot of good commitment research. I think it's unlikely the labs will make much of an attempt to box their AIs.
- In practice, I think we probably go from an AI that reasonably assigns 10% to takeover success to an AI that reasonably assigns 1-eps to takeover success in a single training run / development process and very quickly (like, plausibly in days).

(that said, i still think this plan is good/interesting relative to other existing plans...)